

See discussions, stats, and author profiles for this publication at: <https://www.researchgate.net/publication/238717518>

A Turbo Tutorial

Article · January 1999

CITATIONS

0

READS

803

1 author:



[Jakob Dahl Andersen](#)

Technical University of Denmark

24 PUBLICATIONS **214** CITATIONS

SEE PROFILE

A Turbo Tutorial

by Jakob Dahl Andersen

Department of Telecommunication

Technical University of Denmark

<http://www.tele.dtu.dk/~jda/>

1. Introduction	3
2. Turbo Codes	5
2.1 Encoding	5
2.2 First Decoding	8
2.3 Putting Turbo on the Turbo Codes.	9
2.4 Performance Example	11
3. APP Decoding	13
4. Final Remarks	19
5. Selected Literature	21

1. Introduction

The theory of error correcting codes has presented a large number of code constructions with corresponding decoding algorithms. However, for applications where very strong error correcting capabilities are required these constructions all result in far too complex decoder solutions. The way to combat this is to use concatenated coding, where two (or more) constituent codes are used after each other or in parallel - usually with some kind of interleaving. The constituent codes are decoded with their

respective decoders, but the final decoded result is usually sub-optimal. This means that better results might be achieved with a more complicated decoding algorithm - like the brute-force trying of all possible codewords. However, concatenated coding offers a nice trade of between error correcting capabilities and decoder complexity.




Figure 1 Concatenated coding.

Concatenated coding is illustrated in Figure 1. Here we see the information frame illustrated as a square - assuming block interleaving - and we see the parity from the vertical encoding and the parity from the horizontal encoding. For serial concatenation the parity bits from one of the constituent codes are encoded with the second code and we have parity of parity. If the codes are working in parallel, we do not have this additional parity.

The idea of concatenated coding fits well with Shannon's channel coding theorem, stating that as long as we stay on the right side of the channel capacity we can correct everything - if the code is long enough. This also means that if the code is very long, it does not have to be optimal. The length in itself gives good error correcting capabilities, and concatenated coding is just a way of constructing - and especially decoding - very long codes.

2. Turbo Codes

2.1 Encoding

€

The basic idea of turbo codes is to use two convolutional codes in parallel with some kind of interleaving in between. Convolutional codes can be used to encode a continuous stream of data, but in this case we assume that data is configured in finite blocks - corresponding to the interleaver size.




Figure 2 Turbo encoder

The frames can be terminated - i.e. the encoders are forced to a known state after the information block. The termination tail is then appended to the encoded information and used in the decoder. The system is illustrated in Figure 2.

We can regard the turbo code as a large block code. The performance depends on the weight distribution - not only the minimum distance but the number of words with low weight. Therefore, we want input patterns giving low weight words from the first encoder to be interleaved to patterns giving words with high weight for the second encoder.

Convolutional codes have usually been encoded in their feed-forward form, like $(G1, G2) = (1+D^2, 1+D+D^2)$. However, for these codes a single 1, i.e. the sequence ...0001000..., will give a codeword which is exactly the generator vectors and the weight of this codeword will in general be very low. It is clear that a single 1 will propagate through any interleaver as a single 1, so the conclusion is that if we use the codes in the feed-forward form in the turbo scheme the resulting code will have a large number of codewords with very low weight.

The trick is to use the codes in their recursive systematic form where we divide with one of the generator vectors. Our example gives $(1, G2/G1) = (1, (1+D+D^2)/(1+D^2))$. This operation does not change the set of encoded sequences, but the mapping of input sequences to output sequences is different. We say that the code is the same, meaning that the distance properties are unchanged, but the encoding is different.

In Figure 3 we have shown an encoder on the recursive systematic form. The output sequence we got from the feed-forward encoder with a single 1 is now obtained with the input $1+D^2=G1$. More important is the fact that a single 1 gives a codeword of semi-infinite weight, so with the recursive systematic encoders we may have a chance to find an interleaver where information patterns giving low weight words from the first encoder are interleaved to patterns giving words with high weight from the second encoder. The most critical input patterns are now patterns of weight 2. For the example code the information sequence ...01010... will give an output of weight 5.

Notice that the fact that the codes are systematic is just a coincidence, although it turns out to be very convenient for several reasons. One of these is that the bit error rate (BER) after decoding of a systematic code can not exceed the BER on the channel. Imagine that the received parity symbols were completely random, then the decoder would of course stick to the received version of the information. If

the parity symbols at least make some sense we would gain information on the average and the BER after decoding will be below the BER on the channel.




Figure 3 Recursive systematic encoder

One thing is important concerning the systematic property, though. If we transmit the systematic part from both encoders, this would just be a repetition, and we know that we can construct better codes than repetition codes. The information part should only be transmitted from one of the constituent codes, so if we use constituent codes with rate 1/2 the final rate of the turbo code becomes 1/3. If more redundancy is needed, we must select constituent codes with lower rates. Likewise we can use puncturing after the constituent encoders to increase the rate of the turbo codes.

Now comes the question of the interleaving. A first choice would be a simple block interleaver, i.e. to write by row and read by column. However, two input words of low weight would give some very unfortunate patterns in this interleaver. The pattern is shown in Figure 4 for our example code. We see that this is exactly two times the critical two-input word for the horizontal encoder and two times the critical two-input pattern for the vertical

encoder as well. The result is a code word of low weight (16 for the example code) - not the lowest possible, but since the pattern appears at every position in the interleaver we would have a large number of these words.

This time the trick is to use a pseudo-random interleaver, i.e. to read the information bits to the second encoder in a random (but fixed) order. The pattern from Figure 4 may still appear, but not nearly as often. On the other hand we now have the possibility that a critical two-input pattern is interleaved to another critical two-input pattern.

.
.	.	.	0	0	0	0	0	.	.	.
.	.	.	0	1	0	1	0	.	.	.
.	.	.	0	0	0	0	0	.	.	.
.	.	.	0	1	0	1	0	.	.	.
.	.	.	0	0	0	0	0	.	.	.
.

Figure 4 Critical pattern in block interleaver

The probability that a specific two-input pattern is interleaved to another (or the same) specific two-input pattern is $2/N$, where N is the size of the interleaver. Since the first pattern could appear at any of the N positions in the block, we must expect this unfortunate match to appear 2 times in a pseudo-random interleaver of any length. Still the pseudo random interleaver is superior to the block interleaver, and the pseudo-random interleaving is standard for the turbo codes.

It is possible to find interleavers that are slightly better than the pseudo-random ones, some papers on this topic are included in the literature list.

We will end this section by showing a more detailed drawing of a turbo encoder, Figure 5. Here we see the two recursive systematic encoders, this time for the code $(1, (1+D^4)/(1+D+D^2+D^3+D^4))$. Notice that the systematic bit is removed from one of them. At the input of the constituent encoders we see a switch. This is used to force the encoders to the all-zero state - i.e. to terminate the trellis. The complete incoming frame is kept in a buffer from where it is read out with two different sets of ad

dresses - one for the original sequence and one for the interleaved one. This way output 1 and output 2 correspond to the same frame and can be merged before transmission.

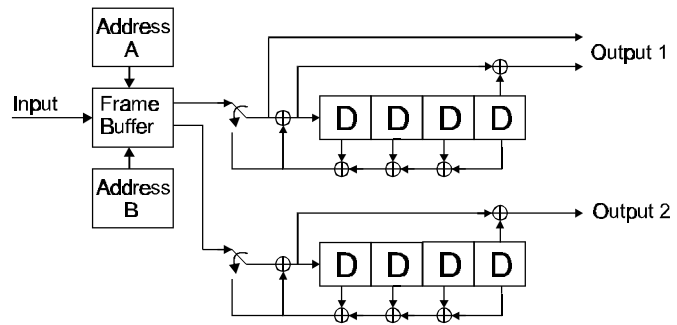


Figure 5 Turbo encoder example

2.2 First Decoding

€

Decoding of error correcting codes is basically a comparison of the probabilities for different codewords - or with convolutional codes, different paths in the trellis. When we talk about probabilities, it is always the probability of some event given a certain amount of information about this event. This is especially clear when we talk about probabilities of something that has already happened - which is always the case in coding theory. What we mean when we talk about the probability that x was sent, $p(x)$, is the probability that x was sent given the amount of information that we have about the event. Usually that is only the received noisy version of x - and of course knowledge of the coding scheme, transmission link etc.

In some cases we have some knowledge of the transmitted signal - before we decode the received one. That may be information that some messages are more likely to occur than others or information from other transmitted sequences. We call this information a priori information and have the corresponding a priori probabilities. Similar we talk about a posteriori probabilities when we have included both the a priori information probabilities and the information gained by the decoding.

For turbo codes we have two encoded sequences. Clearly we must start by decoding one of them to get a first estimate of the information sequence. This estimate should then be used as a priori information in the decoding of the second encoded sequence. This requires that the decoder is able to use a soft decision input and to produce some kind of soft output. The decoding is sketched in Figure 6.




Figure 6 First decoding stage

The standard decoder for turbo codes is the A Posteriori Probability decoding (APP) (sometimes referred to as the Maximum A Posteriori decoding algorithm (MAP)). The APP decoder, described in Section 3, does indeed calculate the a posteriori probabilities for each information bits.

We will represent the soft input/output as log-likelihood ratios, i.e. a signed number where negative numbers indicate that zero is the most likely value of the bit. As seen from Formula 1 the log-likelihood ratio of the a posteriori probabilities can easily be divided into two components - the log-likelihood ratio of the a priori probabilities of the bit d_t and the information gained by the current observation. This means that when we gain additional information about the information bits - like with the second decoding - we simply add a (negative or positive) component to the log-likelihood ratio.

$$\begin{aligned}
 \Lambda(d_t) &= \log \frac{\Pr\{d_t=1, \text{observation}\}}{\Pr\{d_t=0, \text{observation}\}} \\
 &= \log \frac{\Pr_{ap}\{d_t=1\}}{\Pr_{ap}\{d_t=0\}} + \log \frac{\Pr\{\text{observation} | d_t=1\}}{\Pr\{\text{observation} | d_t=0\}} \\
 &= \log \frac{\Pr_{ap}\{d_t=1\}}{\Pr_{ap}\{d_t=0\}} + \Lambda'(d_t)
 \end{aligned} \tag{1}$$

2.3 Putting Turbo on the Turbo Codes.

€

When we have a parity equation, it involves a number of information bits. Let us look at one of the simplest possible parity equations - a sum of two information bits: $P=I_1+I_2$. It is clear that if both P and I_2 are very reliable we get a reliable estimate of I_1 , on the other

hand if I_2 is very unreliable we do not get much information about I_1 . If we now imagine that both I_1 and I_2 are unreliable when we decoded the first sequence, but that I_2 is involved in some parity equations with very reliable bits in the second encoded sequence - then we might return to the parity equations from the first sequence for a second iteration with this new and much more reliable estimate of I_2 . This way we could continue to decode the two encoded sequences and iterate towards the final decision.

However, it is not that easy since we must be very careful not to use our information more than once. Luckily we see from Formula 1, that it is easy to subtract the a priori information - which came from the other decoder - from the decoder output. This will prevent most of the unwanted positive feed-back. We may still have loops in the decision process, though, i.e. we might see that I_1 influences I_2 in the first decoder, that I_2 influences I_3 in the second decoder and finally that I_3 influences I_1 in the next iteration in the first decoder. This way the new improved estimate of I_1 will be based on information that came from I_1 in the first place.

Use of the system in practice has shown that if we subtract the log-likelihood ratio of the a priori information after each constituent decoder and make a number of decoding iterations we get a system that is working remarkably well - for many applications it actually outperforms the previously known systems. Still, we must conclude that the final result after turbo decoding is a sub-optimal decoding due to the loops in the decision process. For low signal-to-noise ratios we may even see that the decoding does not converge to anything close to the transmitted codeword.

The turbo decoder is shown in Figure 7.

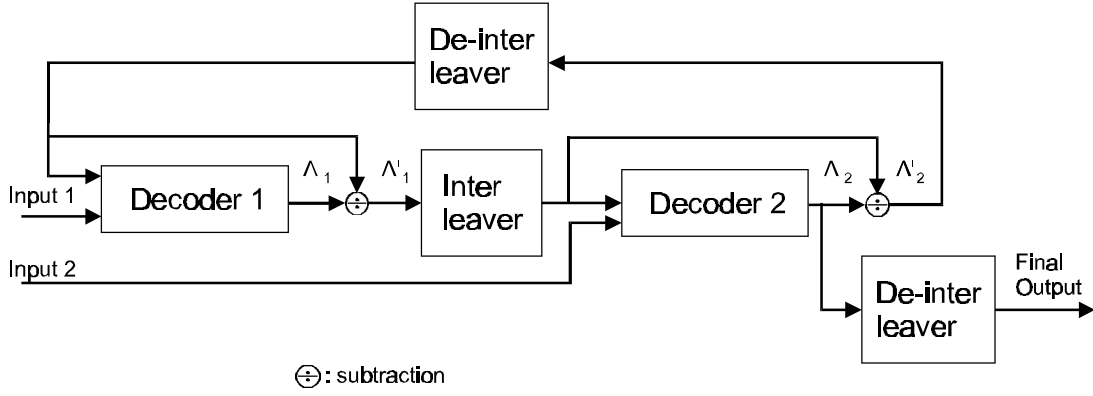


Figure 7 Turbo decoder

2.4 Performance Example

€

We will show an example of the performance with turbo codes. We use the system illustrated in Figure 5, i.e. the code $(1, (1+D^4)/(1+D+D^2+D^3+D^4))$ for both encoders but the information sequence is only transmitted from the first one. This means that the over-all rate is $1/3$. The block length is 10384 bits and we use a pseudo-random interleaver. After each frame the encoders are forced to the zero state. The corresponding termination tail - 4 information bits and 4 parity bits for each encoder, a total of 16 bits - is appended to the transmitted frame and used in the decoder. In principle the termination reduces the rate, but for large frames this has no practical influence. In this case the rate is reduced from 0.3333 to 0.3332.

The performance curves for Bit Error Rate (BER) and Frame Error Rate (FER) are shown in Figure 8. Due to the sub-optimal decoding the performance curves consist of two parts.

For low signal-to-noise ratios the main problem is lack of convergence in the iterated decoding process, resulting in frames with a large number of errors. In this region we are far from optimal decoding. This means that we may benefit from more iterations. As we see from the figure there is a considerable gain by going from 8 to 18 iterations, and with more iterations the performance might be even better.

For high signal-to-noise ratios the decoding is almost optimal, and the main problem is codewords of low weight. This region is usually referred to as the error-floor since the improvement for increasing signal-to-noise ratio is very small. In spite of the name it is not a true floor, since the BER and FER is constantly decreasing - although not nearly as fast as for the low signal-to-noise ratios. Notice that when the signal to noise ratio is high a small number of iterations is sufficient.




Figure 8 Simulation results.

3. APP Decoding

The A Posteriori Probability (APP) algorithm does in fact calculate the *a posteriori* probabilities of the transmitted information bits for a convolutional code. In this presentation we will restrict ourselves to convolutional codes with rate $1/n$.

The convolutional encoder with memory M (Figure 3) may be seen as a Markov source with 2^M states S_t , input d_t and output \mathbf{X}_t . The output \mathbf{X}_t and the new state S_t are functions of the input d_t and the previous state S_{t-1} .

If the output \mathbf{X}_t is transmitted through a Discrete Memoryless Channel with white Gaussian noise. The probability of receiving \mathbf{Y}_t when \mathbf{X}_t was sent is

$$\Pr\{\mathbf{Y}_t | \mathbf{X}_t\} = \prod_{j=1}^n \frac{1}{\sqrt{2\pi}\sigma} \cdot e^{-\frac{(y_{tj} - x_{tj})^2}{2\sigma^2}} \quad (2)$$

where x_{tj} is the j -th bit of the transmitted word \mathbf{X}_t , and y_{tj} the corresponding received value. The signal to noise ratio is $E_s/N_0 = 1/2\sigma^2$. In principle knowledge of the signal-to-noise ratio is needed for the APP algorithm. However, it may be chosen to a fixed value - depending on the operation point of the system - with only a small degradation of the performance.

Assume that we receive the sequence $\mathbf{Y}_1^L = \mathbf{Y}_1, \mathbf{Y}_2, \dots, \mathbf{Y}_L$. The a posteriori probabilities of the state transitions (i.e. branches) are found as:

$$\Pr\{S_{t-1} = m', S_t = m | \mathbf{Y}_1^L\} = \frac{\Pr\{S_{t-1} = m', S_t = m, \mathbf{Y}_1^L\}}{\Pr\{\mathbf{Y}_1^L\}}, \quad t = 1 \dots L \quad (3)$$

$\Pr\{\mathbf{Y}_1^L\}$ is a constant for a given received sequence and since we consider rate $1/n$ codes only, there is one specific information bit associated with each state transition. We therefore define

$$\sigma_t(i, m') = \Pr\{d_t=i, S_{t-1}=m', \mathbf{Y}_1^L\} \quad (4)$$

The final log-likelihood ratio becomes

$$\Lambda(d_t) = \log \frac{\Pr\{d_t=1, \text{observation}\}}{\Pr\{d_t=0, \text{observation}\}} = \frac{\sum_{m'} \sigma_t(1, m')}{\sum_{m'} \sigma_t(0, m')} \quad (5)$$

In order to calculate $\sigma_t(i, m')$ we define the following probability functions $\alpha_t(m)$, $\beta_t(m)$, and $\gamma_t(i, m')$ as

$$\alpha_t(m) = \Pr\{S_t=m, \mathbf{Y}_1^t\} \quad (6)$$

$$\beta_t(m) = \Pr\{\mathbf{Y}_{t+1}^L | S_t=m\} \quad (7)$$

$$\gamma_t(i, m') = \Pr\{d_t=i, \mathbf{Y}_t | S_{t-1}=m'\} \quad (8)$$

Compared to the Viterbi algorithm $\alpha_t(m)$ corresponds to the state metrics, while $\gamma_t(i, m')$ corresponds to the branch metrics. $\beta_t(m)$ can be seen as backwards state metrics.

For the notation we will also need the function giving the new encoder state S_t when $S_{t-1}=m'$ and $d_t=i$

$$\text{newstate}(i, m')$$

and the function giving the old encoder state S_{t-1} when $S_t=m$ and $d_t=i$

$$\text{oldstate}(i, m)$$

Since the encoder is a Markov process and the channel is memoryless, we have

$$\Pr\{\mathbf{Y}_{t+1}^L | S_t=m, \mathbf{Y}_1^t\} = \Pr\{\mathbf{Y}_{t+1}^L | S_t=m\} \quad (11)$$

and

$$\begin{aligned}
\sigma_t(i, m') &= \Pr\{S_{t-1} = m', \mathbf{Y}_1^{t-1}\} \cdot \Pr\{d_t = i, \mathbf{Y}_t \mid S_{t-1} = m'\} \\
&\quad \cdot \Pr\{\mathbf{Y}_{t+1}^L \mid S_t = \text{newstate}(i, m')\} \\
&= \alpha_{t-1}(m') \cdot \gamma_t(i, m') \cdot \beta_t(\text{newstate}(i, m'))
\end{aligned} \tag{12}$$

If we assume that the frames are terminated to state 0, we have $\alpha_0(0)=1$, and $\alpha_0(m)=0$, $m=1,2,\dots,2^M-1$. We can calculate α as a forward recursion

$$\begin{aligned}
\alpha_t(m) &= \sum_{i=0,1} \Pr\{d_t = i, S_{t-1} = \text{oldstate}(i, m), \mathbf{Y}_1^t\} \\
&= \sum_{i=0,1} \Pr\{S_{t-1} = \text{oldstate}(i, m), \mathbf{Y}_1^{t-1}\} \cdot \Pr\{d_t = i, \mathbf{Y}_t \mid S_{t-1} = \text{oldstate}(i, m)\} \\
&= \sum_{i=0,1} \alpha_{t-1}(\text{oldstate}(i, m)) \cdot \gamma_t(i, \text{oldstate}(i, m))
\end{aligned} \tag{13}$$

At the end of the frame we have $\beta_L(0)=1$, and $\beta_L(m)=0$, $m=1,2,\dots,2^M-1$. We can calculate β as a backward recursion

$$\begin{aligned}
\beta_t(m) &= \sum_{i=0,1} \Pr\{d_{t+1} = i, \mathbf{Y}_{t+1}^L \mid S_t = m\} \\
&= \sum_{i=0,1} \Pr\{d_{t+1} = i, \mathbf{Y}_{t+1} \mid S_t = m\} \cdot \Pr\{\mathbf{Y}_{t+2}^L \mid S_{t+1} = \text{newstate}(i, m)\} \\
&= \sum_{i=0,1} \gamma_{t+1}(i, m) \cdot \beta_{t+1}(\text{newstate}(i, m))
\end{aligned} \tag{14}$$

If the frames are not terminated we have no knowledge of the initial and final states. In this case we must use $\alpha_0(m) = \beta_L(m) = 2^{-M}$.

Since $\alpha_t(m) = \Pr\{S_t = m, \mathbf{Y}_1^t\}$ becomes very small with increasing t some rescaling must be used. In principle the function $\alpha'_t(m)$ should be used

$$\alpha'_t(m) = \Pr\{S_t = m \mid \mathbf{Y}_1^t\} = \frac{\Pr\{S_t = m, \mathbf{Y}_1^t\}}{\Pr\{\mathbf{Y}_1^t\}} \tag{15}$$

where $\Pr\{\mathbf{Y}_1^t\}$ is found as the sum of $\alpha_t(m)$ over all states, meaning that the $\alpha'_t(m)$ values always add up to one. However, since the output is the log-likelihood ratio the actual

rescaling is not important as long as underflows are avoided. Similar the function $\beta_t(m)$ needs rescaling.

The algorithm sketched here requires that $\alpha_t(m)$ is stored for the complete frame since we have to await the end of the frame before we can calculate $\beta_t(m)$. We can instead use a sliding window approach with period T and training period Tr . First $\alpha_t(m)$ is calculated and stored for $t=0$ to $T-1$. The calculation of $\beta_t(m)$ is initiated at time $t=T+Tr-1$ with initial conditions $\beta_{T+Tr-1}(m)=2^{-M}$. The first Tr values $\beta_t(m)$ is discarded but after the training period, i.e. for $t=T-1$ down to 0 , we assume that $\beta_t(m)$ is correct and ready for the calculation of $\sigma_t(i, m')$. After the first window we continue with the next one until we reach the end of the frame where we use the true final conditions for $\beta_L(m)$.

Of course, this approach is an approximation but if the training period is carefully chosen the performance degradation can be very small.

Since we have only one output associated with each transition, we can calculate $\gamma_t(i, m')$ as

$$\gamma_t(i, m') = \Pr_{\text{apriori}}\{d_t = i\} \cdot \Pr\{\mathbf{Y}_t \mid d_t = i, S_{t-1} = m'\} \quad (16)$$

For turbo codes the a priori information typically arrives as a log-likelihood ratio. Luckily we see from the calculation of $\alpha_t(m)$ and $\beta_t(m)$ that $\gamma_t(i, m')$ is always used in pairs - $\gamma_t(0, m')$ and $\gamma_t(1, m')$. This means we can multiply $\gamma_t(i, m')$ with a constant

$$k_t = \frac{1}{\Pr_{\text{apriori}}\{d_t = 0\}} \quad (17)$$

and get

$$\gamma'_t(1, m') = \frac{\Pr_{\text{apriori}}\{d_t = 1\}}{\Pr_{\text{apriori}}\{d_t = 0\}} \cdot \Pr\{\mathbf{Y}_t \mid d_t = 1, S_{t-1} = m'\} \quad (18)$$

$$\gamma'_t(0, m') = \Pr\{\mathbf{Y}_t \mid d_t = 0, S_{t-1} = m'\} \quad (19)$$

For an actual implementation the values of $\alpha_t(m)$, $\beta_t(m)$ and $\gamma_t(i,m')$ may be represented as the negative logarithm to the actual probabilities. This is also common practice for Viterbi decoders where the branch and state metrics are $-\log$ to the corresponding probabilities.

With the logarithmic presentation multiplication becomes addition and addition becomes an E-operation, where

$$x \text{ E } y = -\log(e^{-x} + e^{-y}) = \min(x, y) - \log(1 + e^{-|y-x|}) \quad (20)$$

This function can be reduced to finding the minimum and adding a small correction factor.

As seen from Formula 18 the incoming log-likelihood ratio Λ' , can be used directly in the calculation of $-\log(\gamma)$ as the log-likelihood ratio of the a priori probabilities.

4. Final Remarks

This tutorial was meant as a first glimpse on the turbo codes and the iterated decoding principle. Hopefully, we have shed some light on the topic, if there are still some dark spots - try reading it again!

Of course there are a lot of details not explained here, a lot of variation to the turbo coding scheme and a lot of things that may need a proof. Some of these can be found in the papers on the literature list.

5. Selected Literature

- [1] Jakob Dahl Andersen, “Turbo Codes Extended with Outer BCH Code”, *Electronics Letters*, vol. 32 No. 22, Oct. 1996.
- [2] J. Dahl Andersen and V. V. Zyablov, “Interleaver Design for Turbo Coding”, Proc. Int. Symposium on Turbo Codes, Brest, Sept. 1997.
- [3] Jakob Dahl Andersen, “Selection of Component Codes for Turbo Coding based on Convergence Properties”, *Annales des Telecommunication, Special issue on iterated decoding*, June 1999.
- [4] L. R. Bahl, J. Cocke, F. Jelinek and R. Raviv, “Optimal Decoding of Linear Codes for Minimizing Symbol Error Rate,” *IEEE Trans. Inform. Theory*, vol IT-20, pp 284-287, March 1974.
- [5] S. Benedetto and G. Montorsi, “Performance Evaluation of Turbo-codes”, *Electronics Letters*, vol. 31, No. 3, Feb. 1995.
- [6] S. Benedetto and G. Montorsi, “Serial Concatenation of Block And Convolutional Codes”, *Electronics Letters*, Vol. 32, No. 10, May 1996.
- [7] C. Berrou, A. Glavieux and P. Thitimajshima, “Near Shannon Limit Error-correcting Coding and Decoding : Turbo-codes(1)”, Proc. ICC '93, pp. 1064-1070, May 1993.
- [8] C. Berrou and A. Glavieux, “Near Optimum Error Correcting Coding and Decoding: Turbo Codes”, *IEEE trans. on Communications*, Vol. 44, No. 10, Oct. 1996.
- [9] R. J. McEliece, E. R. Rodemich and J.-F. Cheng, “The Turbo Decision Algorithm”, Presented at the 33rd Allerton Conference on Communication, Control and Computing, Oct. 1995.

- [10] L. C. Perez, J. Seghers and D. J. Costello, Jr, "A Distance Spectrum Interpretation of Turbo Codes", *IEEE Trans. on Inform. Theory*, Vol. 42, No. 6, Nov. 1996.
- [11] Steven S. Pietrobon, "Implementation and Performance of a Serial MAP Decoder for use in an Iterative Turbo Decoder", Proc. Int. Symposium on Information Theory, Whistler, Canada, Sept. 1995.